

# Predicting Air BnB Prices

*Andrew Cohen*

*1/10/2020*

## R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
library(readr)
library(ggplot2)
library(plyr)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:plyr':
##
##   arrange, count, desc, failwith, id, mutate, rename, summarise,
##   summarize

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(xgboost)
```

```
##
## Attaching package: 'xgboost'

## The following object is masked from 'package:dplyr':
##
##   slice
```

```
library(dummies)
```

```
## dummies-1.5.6 provided by Decision Patterns
```

```
library(rsample)
```

```
## Loading required package: tidyr
```

```
df_listings <- read_csv("/Users/cohean/Desktop/DataSciChallenge/listings.csv",
                        col_types = cols(host_id = col_character(),
                                         id = col_character()))
```

```
# EDA
# library(rpivotTable)
# rpivotTable(df_listings)
```

```
summary(df_listings)
```

```
##      id          name      host_id
## Length:48864    Length:48864    Length:48864
## Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character
##
##
##
##      host_name      neighbourhood_group neighbourhood      latitude
## Length:48864    Length:48864      Length:48864    Min.   :40.50
## Class :character Class :character      Class :character    1st Qu.:40.69
## Mode  :character Mode  :character      Mode  :character    Median :40.72
##                                           Mean   :40.73
##                                           3rd Qu.:40.76
##                                           Max.   :40.91
##
##      longitude      room_type      minimum_nights
## Min.   :-74.24    Length:48864    Min.    : 1.000
## 1st Qu.: -73.98    Class :character    1st Qu.: 1.000
## Median : -73.96    Mode  :character    Median : 2.000
## Mean   : -73.95                                Mean  : 7.093
## 3rd Qu.: -73.94                                3rd Qu.: 5.000
## Max.   : -73.71                                Max.   :1250.000
##
##      calculated_host_listings_count availability_365 number_of_reviews
## Min.   : 1.000                                Min.   : 0.0    Min.   : 0.00
## 1st Qu.: 1.000                                1st Qu.: 0.0    1st Qu.: 1.00
## Median : 1.000                                Median : 41.0    Median : 5.00
## Mean   : 7.438                                Mean   :112.5    Mean   : 23.39
## 3rd Qu.: 2.000                                3rd Qu.:232.0    3rd Qu.: 24.00
## Max.   :343.000                                Max.   :365.0    Max.   :639.00
##
##      reviews_per_month      price
## Min.   : 0.010    Min.   : 0.0
## 1st Qu.: 0.190    1st Qu.: 69.0
## Median : 0.710    Median : 105.0
## Mean   : 1.366    Mean   : 151.5
## 3rd Qu.: 2.000    3rd Qu.: 175.0
## Max.   :66.610    Max.   :10000.0
## NA's   :10131
```

```
# remove nuisance columns
```

```
# additional analysis thoughts, potentially can include the "name" column using nlp techniques
df_listings <- subset(df_listings, select = -c(id, name, host_id, host_name, neighbourhood))
```

```
# check categorical for errors vars before encoding
```

```
df_listings %>% count(neighbourhood_group, sort = TRUE)
```

```
## # A tibble: 5 x 2
##   neighbourhood_group      n
##   <chr>                <int>
## 1 Manhattan            21456
```

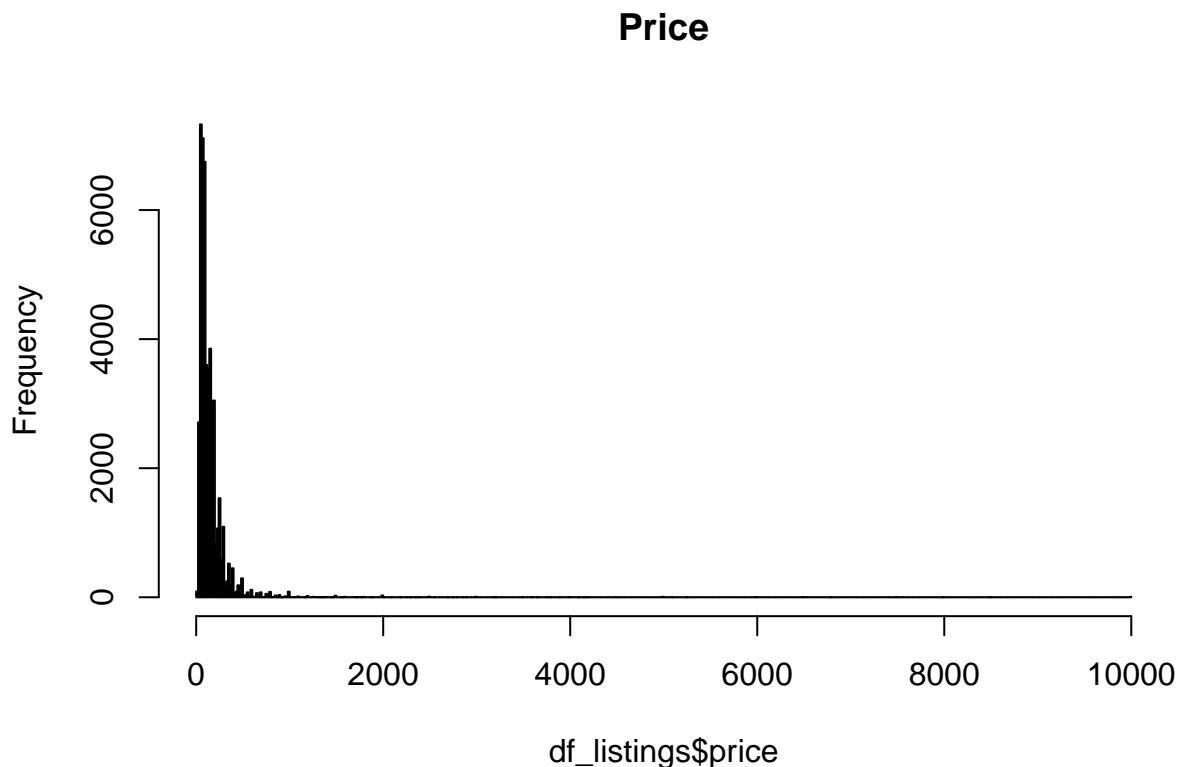
```
## 2 Brooklyn      20114
## 3 Queens        5811
## 4 Bronx         1105
## 5 Staten Island  378
```

```
df_listings %>% count(room_type, sort = TRUE)
```

```
## # A tibble: 3 x 2
##   room_type      n
##   <chr>      <int>
## 1 Entire home/apt 25296
## 2 Private room   22397
## 3 Shared room    1171
```

```
#df_listings %>% count(neighbourhood, sort = TRUE)
```


```
# lapply(df_listings,class)
hist(df_listings$price, breaks = 500, main = "Price")
```




```
# boxplot(df_listings$price)

#remove outliers
df_model <- df_listings[df_listings$price<3500,]
# remove zeros
df_model <- df_model[df_model$price != 0,]
hist(df_model$price, breaks = 250, main = "Price")
```


## Price




```
# boxplot room type  
ggplot(df_model, aes(x=room_type, y=price, fill=room_type)) +  
  geom_boxplot() +  
  coord_cartesian(ylim = c(0, 500)) #zoom into center of data
```



```
# boxplot boro
ggplot(df_model, aes(x=neighbourhood_group, y=price, fill=neighbourhood_group)) +
  geom_boxplot() +
  coord_cartesian(ylim = c(0, 500)) #zoom into center of data
```



```
# grouped boxplot
ggplot(df_listings, aes(x=neighbourhood_group, y=price, fill=room_type)) +
  geom_boxplot() +
  coord_cartesian(ylim = c(0, 500)) #zoom into center of data
```



```
numericV <- which(sapply(df_model, is.numeric))
cor(df_model$price, df_model[, numericV])

##          latitude longitude minimum_nights calculated_host_listings_count
## [1,] 0.04812225 -0.2158065      0.04180992              0.09620459
##      availability_365 number_of_reviews reviews_per_month price
## [1,]      0.1047912      -0.05645767              NA      1
#not much correlation among numeric vars

# impute mean for missing to keep from losing data
df_imp <- transform(df_model,
                    reviews_per_month = ifelse(is.na(reviews_per_month),
                                                mean(reviews_per_month, na.rm=TRUE), reviews_per_month))

# split data for models
df_split <- initial_split(df_imp, prop = .7)
df_train <- training(df_split)
df_test  <- testing(df_split)

# one-hot
df_dummy_train <- dummy.data.frame(df_train, names = c("neighbourhood_group", "room_type"), sep = ".")

## Warning in model.matrix.default(~x - 1, model.frame(~x - 1), contrasts =
## FALSE): non-list contrasts argument ignored

## Warning in model.matrix.default(~x - 1, model.frame(~x - 1), contrasts =
## FALSE): non-list contrasts argument ignored
```

```

df_dummy_test <- dummy.data.frame(df_test, names = c("neighbourhood_group", "room_type") , sep = ".")

## Warning in model.matrix.default(~x - 1, model.frame(~x - 1), contrasts =
## FALSE): non-list contrasts argument ignored

## Warning in model.matrix.default(~x - 1, model.frame(~x - 1), contrasts =
## FALSE): non-list contrasts argument ignored

# seperate X and Y matrices
df_train_x <- subset(df_dummy_train, select = -c(price))
df_train_y <- subset(df_dummy_train, select = c(price))
df_test_x <- subset(df_dummy_test, select = -c(price))
df_test_y <- subset(df_dummy_test, select = c(price))

# using xgboost is one of the best places to start for a predictive model because it usually fits very
xgb.train <- xgb.DMatrix(data = as.matrix(df_train_x), label=as.matrix(df_train_y))
xgb.test <- xgb.DMatrix(data = as.matrix(df_test_x), label=as.matrix(df_test_y))

# parameters based on some light tuning using regression performance metrics like rmse
params <- list(
  booster = "dart",
  #objective = "reg:gamma",
  max.depth = 5,
  eta = 0.007,
  #subsample = 0.60,
  eval_metric = "rmse"
  # ,eval_metric = "mae"
)

xgb.fit<-xgb.train(
  data = xgb.train,
  params = params,
  nrounds = 300, # cut off based on rmse
  #watchlist = list(test=xgb.test, train=xgb.train),
  #verbose = 1
)

# performance
xgb.fit

## ##### xgb.Booster
## raw: 681 Kb
## call:
##   xgb.train(params = params, data = xgb.train, nrounds = 300)
## params (as set within xgb.train):
##   booster = "dart", max_depth = "5", eta = "0.007", eval_metric = "rmse", silent = "1"
## xgb.attributes:
##   niter
## callbacks:
##   cb.print.evaluation(period = print_every_n)
## # of features: 15
## niter: 300
## nfeatures : 15

```




```
# feature importance
xgb.importance(colnames(xgb.train), model = xgb.fit)
```


##		Feature	Gain	Cover	Frequency
## 1:	room_type.Entire home/apt	4.258356e-01	2.000011e-01	0.0337875887	
## 2:	longitude	1.895041e-01	2.502868e-01	0.2046401622	
## 3:	latitude	9.812910e-02	1.190718e-01	0.2021624057	
## 4:	availability_365	8.982579e-02	1.731157e-01	0.1669106881	
## 5:	minimum_nights	7.156341e-02	1.091317e-01	0.1576754139	
## 6:	calculated_host_listings_count	5.877600e-02	2.013640e-02	0.0685888050	
## 7:	number_of_reviews	4.146954e-02	8.326819e-02	0.0925779930	
## 8:	reviews_per_month	1.231771e-02	1.040967e-02	0.0255659421	
## 9:	neighbourhood_group.Manhattan	7.075697e-03	2.442729e-02	0.0130645343	
## 10:	neighbourhood_group.Queens	4.096104e-03	2.864378e-05	0.0180200473	
## 11:	room_type.Private room	8.382327e-04	1.010487e-02	0.0126140331	
## 12:	neighbourhood_group.Bronx	5.639152e-04	1.873163e-06	0.0042797612	
## 13:	neighbourhood_group.Brooklyn	4.828745e-06	1.603896e-05	0.0001126253	

```
# pred price distrubution comparison
xgb.pred <- predict(xgb.fit,xgb.test,reshape=T)
hist(xgb.pred, breaks = 100,main = "XGB Pred Price",xlim = c(-20 , +500))
```

## XGB Pred Price



```
hist(df_test_y$price, breaks = 1000,main = "Price",xlim = c(-20 , +500))
```



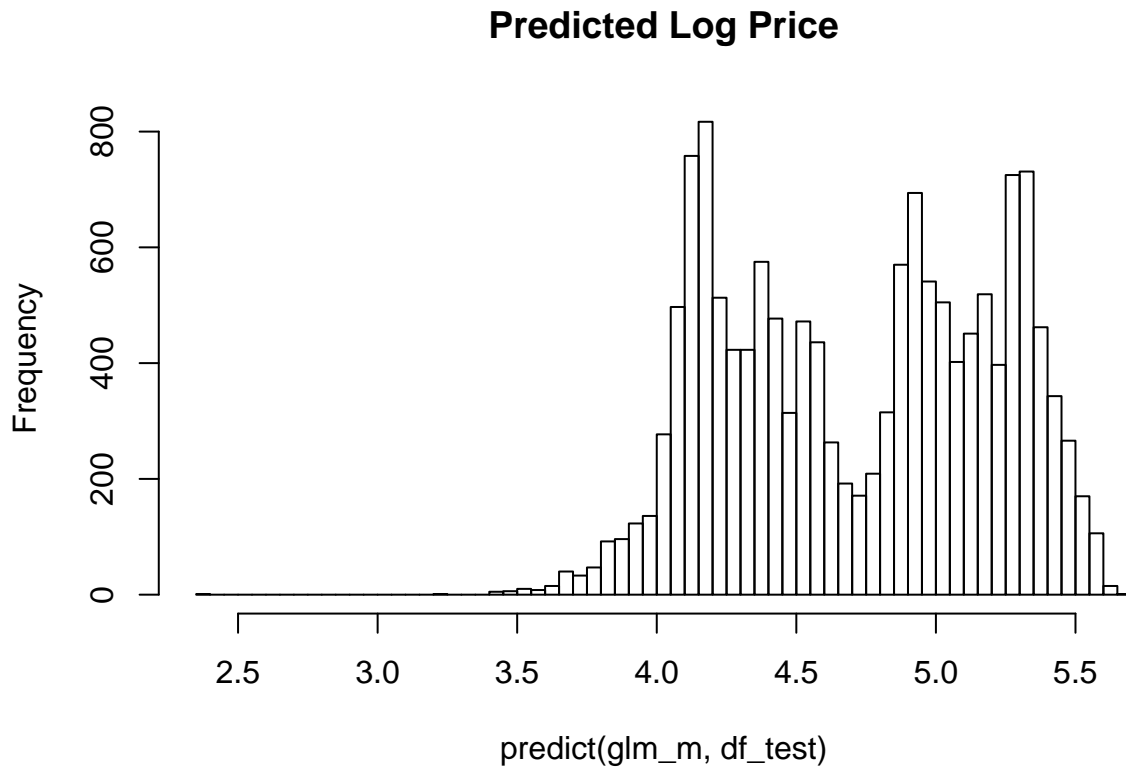
```
# generalized linear model
df_train$logprice <- log(df_train$price)
df_test$logprice <- log(df_test$price)

glm_m <- glm(logprice ~ ., data = subset(df_train, select = -c(price)))
summary(glm_m)
```


```
##
## Call:
## glm(formula = logprice ~ ., data = subset(df_train, select = -c(price)))
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.0250  -0.3096  -0.0483   0.2364   4.2084
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -2.021e+02  8.168e+00  -24.746 < 2e-16
## neighbourhood_groupBrooklyn  -3.144e-02  2.210e-02  -1.422  0.155
## neighbourhood_groupManhattan  2.647e-01  2.001e-02  13.228 < 2e-16
## neighbourhood_groupQueens     8.832e-02  2.113e-02   4.179 2.93e-05
## neighbourhood_groupStaten Island -8.078e-01  4.239e-02 -19.055 < 2e-16
## latitude        -5.595e-01  7.944e-02  -7.044 1.91e-12
## longitude       -3.108e+00  9.188e-02 -33.824 < 2e-16
## room_typePrivate room    -7.469e-01  5.492e-03 -135.996 < 2e-16
## room_typeShared room    -1.153e+00  1.732e-02  -66.538 < 2e-16
## minimum_nights    -2.059e-03  1.355e-04 -15.190 < 2e-16
## calculated_host_listings_count -1.273e-04  8.017e-05  -1.587  0.112
## availability_365     6.915e-04  2.152e-05  32.124 < 2e-16
## number_of_reviews    -8.263e-04  7.135e-05 -11.581 < 2e-16
```

```
## reviews_per_month          1.178e-02  2.184e-03    5.393 6.98e-08
##
## (Intercept)                ***
## neighbourhood_groupBrooklyn
## neighbourhood_groupManhattan ***
## neighbourhood_groupQueens   ***
## neighbourhood_groupStaten Island ***
## latitude                   ***
## longitude                   ***
## room_typePrivate room       ***
## room_typeShared room        ***
## minimum_nights              ***
## calculated_host_listings_count
## availability_365             ***
## number_of_reviews           ***
## reviews_per_month           ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.2348707)
##
##    Null deviance: 16043.0  on 34166  degrees of freedom
## Residual deviance:  8021.5  on 34153  degrees of freedom
## AIC: 47479
##
## Number of Fisher Scoring iterations: 2
```

```
hist(predict(glm_m,df_test), breaks = 100,main = "Predicted Log Price")
```



```
hist(df_test$logprice, breaks = 100, main = "Log Price")
```



```
# both models are pretty good starts  
# xgboost seems better based on the distribution of the residuals  
# similar feature importance
```